

The Information-Saturated Bounce: A De Sitter Attractor from an Assumed Entanglement-Decay Ansatz

Paul Jarvis^{1,*}

¹*Independent Researcher, United Kingdom*

(Dated: July 17, 2026)

We present a coordinate-free toy mechanism for a cyclic cosmological bounce, developed under a specific, stated ansatz for a pre-geometric substrate: that mutual information between relational degrees of freedom decays with distance as a power law, parameterized by a decay exponent α . Given that ansatz, the spectral dynamics of the resulting mutual-information (MI) Laplacian determine the mechanism analyzed below; the ansatz itself is assumed, not derived (Sec. II). By tracking the evolution of the network substrate as it approaches an information-theoretic singularity, we demonstrate that the collapse of spatial locality corresponds to a strict thermodynamic phase transition where the entanglement decay exponent $\alpha \rightarrow 0^+$. Deriving the effective macroscopic energy density from the canonical partition function of the network spectrum, we show that this maximum saturation limit forces the effective equation of state parameter to converge smoothly and continuously onto the dark energy threshold ($w_{\text{eff}} = -1.0$). Simultaneously, the internal thermodynamic energy density saturates at a hard algebraic ceiling fixed by the total number of lattice nodes ($\rho_{\text{eff}} = N$). This invariant cutoff eliminates unphysical density divergences ($\rho \rightarrow \infty$), transforming the attractive cosmic crunch into an inherently regularized, gravitationally repulsive de Sitter expansion phase. We demonstrate the structural stability of this fixed-point attractor across disparate non-spatial graph families under extreme multi-channel chaotic noise. Finally, by mapping the Renormalization Group (RG) flow trajectories and numerically extracting their negative β -functions ($\beta < 0$), we reveal a universal low-dimensional attractor manifold. These results unify microscopic information bounds with cyclic macro-cosmology, consistent with Cosmic Inflation being a structural consequence of an information-saturated quantum crunch within this toy model. We flag three caveats directly rather than leaving them implicit: the equation-of-state identity is exact only in the $\beta \rightarrow 0^+$ limit and is quantified to deviate by $\sim 0.8\%$ at the finite $\beta = 0.1$ used throughout this paper's figures (Table I); our numerical sweeps reproduce this identity by construction from the same scaling relation it is derived from, which confirms the simulation code is correct but is not on its own independent evidence for the underlying physical mechanism; and a previously reported pair of "universal" critical exponents ($A = 0.0166$, $B = 2.5310$) is withdrawn after a finite-size sweep showed both to be artifacts of a single (N , topology) configuration rather than universal quantities, and replaced with a genuinely size-stable exponent $p = 1.122 \pm 0.008$.

extracted by a more robust pooled fit (Sec. V C).

I. INTRODUCTION

A foundational challenge in modern theoretical cosmology is resolving the unphysical singularity that occurs at the boundary of a collapsing universe. In standard General Relativity, a cosmic crunch compresses matter into a zero-volume point, forcing the macroscopic energy density to diverge toward infinity ($\rho \rightarrow \infty$). To circumvent this breakdown of classical spacetime, various frameworks - ranging from Loop Quantum Cosmology to string-inspired braneworlds - postulate ad-hoc fields, high-curvature counter-terms, or external bouncing actions.

In this work, we present a self-contained, coordinate-free alternative, developed under a specific, stated ansatz (Sec. II): that mutual information between relational degrees of freedom decays with distance as a power law. Given that ansatz, the cosmic bounce follows as a direct algebraic consequence of the ansatz's own $\alpha \rightarrow 0$ limit, not as an independently derived physical necessity - we return to this distinction explicitly once the limit itself is introduced below, and again in Sec. IV. Building on recent models where spacetime geometry emerges from the spectrum of mutual-information (MI) Laplacians, we model a cosmic contraction as an information-theoretic compression phase.

As the available spatial volume shrinks, the microscopic degrees of freedom pack tightly together until they saturate the holographic Bekenstein bound. At this absolute saturation threshold, spatial locality undergoes a fundamental phase transition: the entanglement decay exponent collapses ($\alpha \rightarrow 0^+$), transforming the local manifold into a maximally entangled complete graph (K_N). This convergence to K_N is a definitional feature of the power-law ansatz itself - any kernel of the form $I_{ij} \sim d(i, j)^{-\alpha}$ becomes identically distance-independent at $\alpha = 0$ by construction, for any starting topology - not an independently derived physical result. We do not present it as evidence that a real physical system must reach this limit, only as the behavior of this specific model once it does.

Using exact numerical diagonalization on translationally invariant toroidal and non-spatial topological graphs, we map the continuous thermodynamic trajectory of this saturation phase. We demonstrate that as $\alpha \rightarrow 0^+$, the effective cosmological fluid transitions smoothly into a pure dark energy state with an equation of state $w_{\text{eff}} = -1.0$. Concurrently, the emergent energy density approaches a strict algebraic ceiling locked at the total number of system degrees of freedom ($\rho_{\text{eff}} = N$). This dual mechanism provides an organic shock absorber: the invariant energy cap halts the gravitational collapse, while the instantaneous transition to a negative-pressure de Sitter state supplies the overwhelming repulsive force required to fire the universe outward into a new inflationary cycle.

II. RELATIONAL SUBSTRATE AND INVARIANT SCALING MATRIX

We consider a discrete quantum system consisting of N sites with no a priori background geometry. Given a many-body state $|\Psi\rangle$, the network architecture is governed entirely by the

* mrpaulwjarvis@gmail.com

pairwise mutual information $I_{ij} = S(\rho_i) + S(\rho_j) - S(\rho_{ij})$, where S denotes the von Neumann entropy. This relational configuration defines a weighted complete graph whose structural properties are captured by the combinatorial MI-Laplacian matrix:

$$L_{ij}^{\text{MI}} = \delta_{ij} \sum_k I_{ik} - I_{ij}. \quad (1)$$

The eigenvalues λ_n of L^{MI} define the spectral geometry of the emergent manifold.

To strip the framework of any background-dependent mapping or ad-hoc field definitions, the macroscopic energy density ρ_{eff} is derived directly from first-principles statistical mechanics. By treating the eigenvalues of the MI-Laplacian as the foundational spectrum of the relational quantum network's structural fluctuations, we define the canonical partition function over the excited, non-zero spectral states ($\lambda_n > 0$) as:

$$Z = \sum_{n=1}^{N-1} e^{-\beta \lambda_n}, \quad (2)$$

where β is an inverse temperature parameter capturing the thermodynamic scale of the substrate. The internal energy density of the spacetime substrate is then analytically extracted via the standard relation:

$$\rho_{\text{eff}} = -\frac{\partial \ln Z}{\partial \beta} = \frac{\sum_{n=1}^{N-1} \lambda_n e^{-\beta \lambda_n}}{\sum_{n=1}^{N-1} e^{-\beta \lambda_n}}. \quad (3)$$

The application of canonical Boltzmann weights to the spectral eigenvalues is robustly motivated by the algebraic structure of modular Hamiltonians in quantum information theory, rather than an ad-hoc reliance on an external thermal bath. For any localized quantum many-body system, restricting the global vacuum state to a bounded spatial subregion yields an entangled mixed state governed by Tomita–Takesaki modular operators. Under this formulation, the local density matrix takes the exact canonical form $\rho \propto e^{-K}$, where K is the modular Hamiltonian tracking the internal entangling energy. Because the MI-Laplacian L^{MI} functions as the discrete metric proxy for this modular operator, its eigenvalues λ_n map directly onto the spectrum of the system's foundational informational fluctuations. The inverse parameter β does not measure standard thermodynamic temperature, but functions strictly as a microscopic entanglement temperature, setting the characteristic renormalization group (RG) coarse-graining energy scale. At this scale, low-frequency geometric modes are resolved by the relational walk, establishing the partition function Z as a rigorous, observer-independent measure of the substrate's informational capacity.

We parameterize the structural compression of the universe by sweeping the entanglement decay exponent α , which controls the transition from a localized, classical manifold down to a highly non-local saturated network via the power-law correlation profile $I_{ij} \sim d(i, j)^{-\alpha}$, where $d(i, j)$ measures coordinate distances on an $L \times L$ toroidal grid ($N = L^2$). Under an infinitesimal, global volume deformation of the scale factor $d(i, j) \rightarrow a \cdot d(i, j)$, the strict linearity of the combinatorial matrix forces the entire Laplacian to scale uniformly: $L^{\text{MI}}(a) = a^{-\alpha} L^{\text{MI}}(1)$.

Factoring this homogeneous eigenvalue rescaling into the thermodynamic internal energy for-

mulation, the system obeys the classic cosmic fluid relation $\rho_{\text{eff}}(a) \propto a^{-3(1+w_{\text{eff}})}$. Evaluating this ratio analytically yields the definitive centerpiece equation of our model:

$$w_{\text{eff}} = \left(\frac{\alpha}{3}\right) - 1.0. \quad (4)$$

This establishes an exact, continuous mapping between network connectivity and cosmic fluid dynamics. In the highly localized classical limit ($\alpha \rightarrow 4.0$), the substrate naturally recovers the standard radiation-dominated epoch ($w_{\text{eff}} \rightarrow +1/3$). Conversely, as structural compression forces the system toward absolute holographic information saturation, the decay exponent collapses ($\alpha \rightarrow 0^+$), causing the effective equation of state to converge continuously onto the dark energy threshold ($w_{\text{eff}} \rightarrow -1.0$).

III. NUMERICAL VERIFICATION OF THE BOUNCE MECHANISM

To trace this mechanism continuously across the contraction timeline, we execute exact numerical sweeps on an $N = 64$ toroidal lattice, evaluating the network invariants from $\alpha = 1.5$ down to absolute saturation at $\alpha = 0.0$. We note at the outset what this numerical sweep does and does not establish: $w_{\text{eff}} = \alpha/3 - 1.0$ (Sec. II) is an exact analytic identity in the $\beta \rightarrow 0^+$ limit, so a simulation built directly from that same scaling relation reproducing it numerically is an internal consistency check of the simulation code, not independent evidence for the identity itself. The genuinely non-trivial question - how well this identity survives at the finite $\beta = 0.1$ actually used to generate every figure in this paper, where it is only an approximation - is quantified explicitly below.




FIG. 1. The continuous tracking of the effective equation of state $w_{\text{eff}} = \alpha/3 - 1.0$ across the inverted contraction timeline ($\alpha \rightarrow 0^+$). At the absolute information saturation limit ($\alpha = 0$), the substrate fluid converges smoothly and exactly onto the pure cosmological constant threshold ($w_{\text{eff}} = -1.0$).

As demonstrated in Fig. 1, the tracking curve reveals an unbroken, deterministic trajectory. As the universe contracts, the non-local entanglement links expand in capacity, causing α to decrease. Rather than showing chaotic numerical instability or divergence, the effective fluid parameter smoothly plunges down, landing exactly on the dark energy baseline ($w_{\text{eff}} = -1.0000$) at the zero coordinate.

Physically, when $\alpha = 0$, the distance metric becomes structurally irrelevant since $d(i, j)^0 = 1$. The network shifts into a complete graph where every node is globally and identically connected. Because spatial rescalings can no longer dilute or modify the matrix elements, the energy density freezes completely ($\rho_{\text{eff}} \propto a^0$). This constant density sources an instantaneous de Sitter vacuum state characterized by massive negative pressure ($P_{\text{eff}} = -\rho_{\text{eff}}$).

Concurrently, the unphysical infinite density singularities of classical General Relativity are prevented by the network's finite informational storage capacity. As illustrated in Fig. 2, as $\alpha \rightarrow 0^+$, the derived thermodynamic internal energy density climbs before locking horizontally onto a strict algebraic ceiling at exactly $\rho_{\text{eff}} = 64$.




FIG. 2. Numerical evolution of the thermodynamic energy density ρ_{eff} derived from the canonical partition function during spatial contraction. As the correlation decay collapses ($\alpha \rightarrow 0^+$), the internal energy density approaches a hard upper ceiling bounded exactly by the total number of available lattice nodes ($\rho_{\text{eff}} = N = 64$).

This rigid ceiling is an exact algebraic property of complete graphs (K_N) inside the canonical partition function. When $\alpha = 0$, all $N - 1$ excited eigenvalues become perfectly degenerate and lock onto $\lambda_n = N$. Factoring this flat spectrum into Eq. 3 causes the temperature-scaling exponential weights to completely cancel, identically reducing the thermodynamic state parameter to $\rho_{\text{eff}} = N$, exactly and independently of β (Table I). The lattice acts as a cosmological shock absorber: it sets a definitive upper limit on informational compaction, halting the gravitational crunch before a density divergence can occur.

A. Quantifying the finite- β approximation

Away from the $\alpha = 0$ ceiling, the identity $w_{\text{eff}} = \alpha/3 - 1.0$ is exact only as $\beta \rightarrow 0^+$; at finite β the partition weights $e^{-\beta\lambda_n(a)}$ acquire their own residual a -dependence through $\lambda_n(a)$, so the identity becomes an approximation. Table I quantifies this directly at a reference point ($\alpha = 1$, exact $w_{\text{eff}} = -2/3$) across five orders of magnitude in β , using the same $L = 8$ substrate as Figs. 1-2 (`code/beta_invariance_check.py` in the accompanying repository).

β	$\rho_{\text{eff}}(\alpha = 0)$	$w_{\text{eff}}(\alpha = 1)$	$ \Delta w_{\text{eff}} $ (relative)
0.001	64.0000	-0.6667	0.005%
0.01	64.0000	-0.6670	0.056%
0.1	64.0000	-0.6718	0.774%
1.0	64.0000	-0.6947	4.206%
10.0	64.0000	-0.6667	0.000%

TABLE I. β -invariance of the equation-of-state identity. The $\alpha = 0$ ceiling is exact for every β tested, as expected since the spectrum is fully degenerate there and the average is trivially β -independent. The $\alpha = 1$ identity deviates by under 1% (relative) for $\beta \lesssim 0.1$, growing to $\sim 4\%$ at $\beta = 1.0$, and returning to a near-zero value at $\beta = 10$ - a coincidence of this particular test point, not evidence of monotonic behavior. This motivates, but does not by itself fully justify, the use of $\beta = 0.1$ throughout this paper's figures; a systematic study of the β -dependence of every reported quantity is left for future work.

IV. TOPOLOGY INDEPENDENCE OF THE COMPLETE-GRAPH LIMIT

To confirm the numerical implementation behaves correctly for genuinely different starting topologies, not just the toroidal grid used above, we execute a multi-topology finite-size scaling analysis up to $N = 256$ degrees of freedom, across three structurally distinct, non-spatial graph families: Erdős–Rényi random graphs, Watts–Strogatz small-world networks, and Barabási–Albert scale-free architectures.

As illustrated in Fig. 3, regardless of the initial wiring configurations, degree distributions, or clustering coefficients, the absolute information-saturation limit ($\alpha \rightarrow 0^+$) uniformly collapses all coordinate distance dependencies. The substrate identically maps onto a permutation-symmetric complete graph K_N . Across all macro-topologies and scales, the structural spectral gap satisfies $\lambda_1/N = 1.0000$ to machine precision, while the extracted equation of state reveals an absolute null deviation from the dark energy threshold ($|w_{\text{eff}} - (-1.0)| = 0.0000\text{e} + 00$). This equivalence is guaranteed by the ansatz's own functional form: at $\alpha = 0$, every pairwise kernel $I_{ij} \sim d(i, j)^{-\alpha}$ becomes identically 1 regardless of $d(i, j)$, so it no longer matters which distances the starting topology assigned. No other numerical outcome was possible for any of the three families tested.

This is a self-consistency check, not an independent physical result. Because every graph converges to the same complete graph K_N at $\alpha = 0$ by the ansatz's own construction, no outcome other than topology-independence at that limit was possible, regardless of which families were tested - so this does not, on its own, establish that the de Sitter bounce is a property of "any information-saturated quantum substrate" in general. What it does establish is narrower but still useful: that the code correctly implements this guaranteed limit across genuinely different starting topologies,




FIG. 3. Topology-independence self-consistency check evaluated at the absolute saturation boundary ($\alpha = 0$). All three graph families lock onto the same horizontal baseline, as the ansatz guarantees they must.

ruling out an implementation bug that happened to only work for one graph family. Whether a real physical system's mutual-information structure actually follows this power-law ansatz as it approaches saturation - the premise the whole mechanism rests on - is not tested here and remains open.

V. RENORMALIZATION GROUP FLOWS AND ATTRACTOR DYNAMICS

A. Unified 3D Renormalization Group Manifold

The full multivariable synthesis of the cosmological framework is captured by mapping the dynamic trajectories within the unified 3D Renormalization Group (RG) phase space $(\lambda_1, w_{\text{eff}}, d_s)$, tracking the simultaneous evolution of the substrate energy scale, the effective fluid parameter, and the continuous spectral dimension (d_s) .

As illustrated in the 3D phase portrait of Fig. 4, the individual trajectories from widely separated initial network configurations exhibit a powerful low-dimensional convergence. Although the random, small-world, and scale-free networks originate in entirely different sectors of the ultraviolet phase space, the non-linear flow rapidly forces them onto a singular, cohesive attractor manifold. In the infrared limit approaching maximum information saturation ($\lambda_1 \rightarrow N = 100$), all geometric memory of the initial wiring configurations is completely erased, collapsing onto a unique fixed point defined by $\lambda_1 = 100$, $w_{\text{eff}} = -1.0$, and $d_s \approx 5.25$ - the last of these an $N = 100$ -specific value, not a universal constant; it grows systematically with system size (Sec. V C).




FIG. 4. The Universal 3D Renormalization Group Manifold Flow tracking the multi-topology collapse from different structural microstates onto a unified scalar attractor trajectory.

B. Numerical Extraction of the RG β -Functions

To map the explicit differential equations governing the substrate's thermodynamic trajectory, we numerically compute the Renormalization Group (RG) β -functions for both the cosmological fluid state and the emergent spectral geometry with respect to the log-energy scale:

$$\beta_w = \frac{dw_{\text{eff}}}{d \ln \lambda_1}, \quad \beta_{d_s} = \frac{dd_s}{d \ln \lambda_1}. \quad (5)$$

The numerical derivatives are extracted via a centered finite-difference stencil evaluated across the continuous contraction history.

As illustrated in the dual-panel diagnostic of Fig. 5, the extracted β -functions reveal a remarkably consistent structural stability profile. Across all three graph families, both β_w and β_{d_s} remain strictly negative ($\beta < 0$) throughout the entire scaling history, terminating in a tightly focused convergence bundle ($\beta_w \approx -0.30$ and $\beta_{d_s} \approx -0.22$) at the maximum energy scale ($\ln \lambda_1 \rightarrow \ln 100 \approx 4.6$).




FIG. 5. Numerical Renormalization Group β -Functions β_w (left) and β_{d_s} (right). The universal negative sign ($\beta < 0$) across topologies highlights the formal structural stability of the fixed-point attractor.

C. Linearized Stability Criteria and Critical Exponents

To quantify the exact velocity at which the emergent cosmological parameters approach their respective infrared fixed points, we isolate the linearized stability data in the immediate neighborhood of the information-saturation threshold ($\ln \lambda_1 \rightarrow \ln N$). By applying a localized Savitzky–Golay smoothing filter to suppress discrete grid-rearrangement noise, we fit the continuous β -functions to first-order expansion models:

$$\beta_w \approx -A(w_{\text{eff}} + 1), \quad \beta_{d_s} \approx -B(d_s - d_s^*). \quad (6)$$

Fig. 6 shows this procedure applied at a single configuration ($N = 100$, Barabási–Albert topology), yielding $A = 0.0166$ and $B = 2.5310$.




FIG. 6. Smoothed Renormalization Group stability profiles at a single configuration ($N = 100$, Barabási–Albert), displaying linear tracking parameters for the effective fluid scaling exponent A (left) and geometric scale parameter B (right). As shown below, neither value is stable against changes in N or topology.

We tested whether A and B are actually universal by repeating the identical procedure across $N = 50, 100, 200, 400$ and all three graph topologies (`code/critical_exponent_finite_size_scaling.py`). They are not: A ranges over more than two orders of magnitude (from 0.0006 to 0.51) and B changes sign depending on topology – from -7.65 for Erdős–Rényi at $N = 50$

to +5.81 for Barabási–Albert at the same N . The originally quoted pair is an artifact of the single configuration it was measured on, not a critical exponent, and we withdraw the “universal” characterization together with the specific values $A = 0.0166$, $B = 2.5310$.

The same finite-size sweep shows that the spectral-dimension plateau d_s^* is not universal either: it grows linearly with $\ln N$ ($d_s^* = 4.26, 5.25, 6.28, 7.36$ at $N = 50, 100, 200, 400$; regression against $\ln N$ gives slope 1.49 ± 0.02 , $R^2 = 0.9995$). The previously reported $d_s^* \approx 5.25$ is simply the $N = 100$ instance of this size-dependent quantity – the same pattern already found and corrected for the analogous claim in the companion paper *emergent-topological-inflation*, which builds on the same underlying substrate construction.

A genuinely stable quantity does exist in this flow, however. Following the more robust pooled fit used to correct the same issue in that companion paper, we fit the near-boundary flow directly to the power law

$$1 + w_{\text{eff}} = C \left(1 - \frac{\lambda_1}{N} \right)^p, \quad (7)$$

pooling all three topologies at each N . This gives $p = 1.14, 1.12, 1.12, 1.11$ at $N = 50, 100, 200, 400$ respectively (pooled: $p = 1.122 \pm 0.008$, $R^2 = 0.93$), stable to within a few percent across an eightfold change in N – consistent with, and independently reproducing, the exponent $p = 1.11 \pm 0.02$ reported for the same near-boundary attractor behavior in *emergent-topological-inflation*. We report this pooled exponent as the reproducible replacement for the withdrawn A/B pair.




FIG. 7. Pooled near-boundary power-law fit $1 + w_{\text{eff}} = C(1 - \lambda_1/N)^p$ across $N = 50$ – 400 and three graph topologies, giving the size-stable exponent $p = 1.122 \pm 0.008$.

The positive, size-stable exponent p indicates that perturbations away from the de Sitter state are damped as the flow approaches the fixed point – a more modest claim than the originally

asserted exponential damping guaranteed by two independent, strictly positive linear-response coefficients, but one that is actually reproducible. This still supports the non-singular cosmic bounce as a genuine attractor of the scaling flow, rather than an artifact of a single simulated configuration.

VI. DISCUSSION AND CONCLUDING REMARKS

The convergence profiles displayed across our multi-graph scaling engines complete the physical loop required for a fully realized cyclic cosmology. When a collapsing universe hits the information-saturation threshold:

1. Spatial locality dissolves entirely as the coordinates contract, and the graph unifies into a fractional spectral dimension that plateaus near $d_s^* \approx 5.25$ at this paper's reference size ($N = 100$) – a size-dependent value rather than a universal constant (Sec. V C).
2. The energy density hits a hard, built-in structural cutoff at $\rho_{\text{eff}} = N$, absorbing the crunch.
3. The simultaneous transition to a $w_{\text{eff}} = -1.0$ dark energy state instantly converts the attractive gravitational force into an overwhelming cosmological repulsion.

This combination of an energy ceiling and gravitational repulsion acts as an organic bounce mechanism within this toy model's ansatz. The repulsive vacuum pressure forces the system into an explosive, outward-bound trajectory, initiating Cosmic Inflation as a structural consequence of the assumed entanglement-decay ansatz reaching its $\alpha \rightarrow 0$ limit, not as an independently derived physical necessity, seamlessly resetting the universe for its next expansion cycle.

ACKNOWLEDGMENTS

The author thanks the open-source scientific python community for providing the matrix initialization and linear algebra diagnostic suites that enabled the numerical validation of this framework.

-
- [1] A. Connes and M. Marcolli, *Noncommutative Geometry, Quantum Fields and Motives*, American Mathematical Society (2008).
 - [2] A. H. Chamseddine and A. Connes, *The Spectral Action Principle*, Commun. Math. Phys. 186, 731-750 (1997).
 - [3] A. Ashtekar, T. Pawłowski, and P. Singh, *Quantum Nature of the Big Bang: Improved dynamics for loop quantum cosmology*, Phys. Rev. D 74, 084003 (2006).
 - [4] C. J. Cao, S. M. Carroll, and S. Michalakis, *Space from Hilbert Space: Recovering geometry from intrinsic quantum mechanics*, Phys. Rev. D 95, 024031 (2017).
 - [5] T. Konopka, F. Markopoulou, and S. Severini, *Quantum Graphity: A model of emergent locality*, Phys. Rev. D 77, 104029 (2008).